



Хэрэглээний шинжлэх ухаан, Инженерчлэлийн сургууль  
Мэдээлэл, компьютерийн ухааны тэнхим

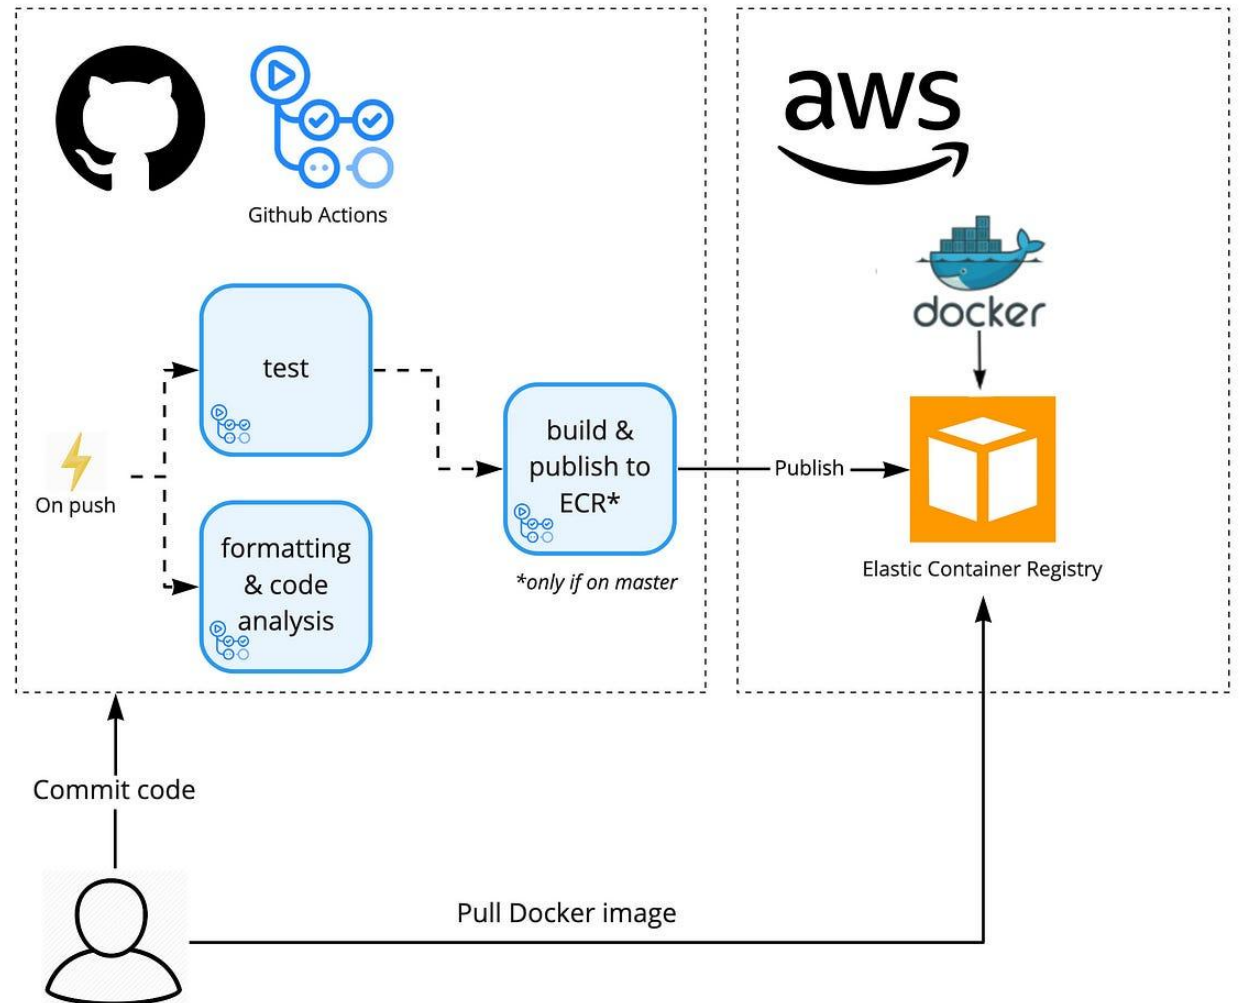
# Week 4

# End-to-End Deployment Workshop

Advanced Web Application Development  
National University of Mongolia | Fall 2025  
Lecturer R.Javkhlan

# Today's Goals

- A live HTTPS URL serving Yellow Books (Next.js web + API).
- CI/CD pipeline: GitHub Actions → ECR → EKS (no static AWS keys).
- Kubernetes resources: Deployments, Services, Ingress (ALB), Secrets/ConfigMaps, Migration Job, and HPA.
- A repeatable DEPLOY.md you can reuse for future projects.





# Reading Map (Modules & Time)

- Module 0: Setup & Repo
- Module 1 — Docker & ECR + GitHub Actions (Build & Push Now)
- Module 2: Kubernetes & EKS
- Module 3: Ingress, TLS, DNS
- Module 4: IAM, OIDC, IRSA
- Module 5: K8s Manifests + Prisma Job
- Module 6: Verify, Troubleshoot, Teardown



Хэрэглээний шинжлэх ухаан, Инженерчлэлийн сургууль  
Мэдээлэл, компьютерийн ухааны тэнхим

# Module 0 — Repo & Local Hygiene



# Create/Prepare GitHub Repo

1. Create a private repo (name idea: yellowbooks-eks).
2. Initialize git locally, commit, and push to main.
3. Set up SSH or HTTPS auth (SSH recommended).
4. Turn on branch protection for main.
5. Sanity-check with two commands.

# Steps. Create repo on GitHub

1. Web UI: GitHub → New → Private → Name:  
yellowbooks-eks → Add README ✓ → Create.
2. Or CLI:

```
# from your project folder (containing apps/, libs/, prisma/ ...)  
gh repo create <ORG/REPO> --private --source=. --remote=origin --push
```

3. Configure git (first time only)

```
git config --global user.name "Your Name"  
git config --global user.email "you@example.com"
```

# Steps. Add remote & push main

```
# If repo was not created with gh CLI
git init -b main
git add .
git commit -m "chore: initial commit (apps/web, apps/api, libs, prisma,
# choose ONE remote style
# SSH (recommended)
git remote add origin git@github.com:<ORG/REP0>.git
# or HTTPS
git remote add origin https://github.com/<ORG/REP0>.git

git push -u origin main
```



# Steps. Enable branch protection (UI)

Settings → Branches → Branch protection rules → Add rule:

- Branch name pattern: main
- ✓ Require a pull request before merging (optional for class)
- ✓ Require status checks to pass before merging (we'll wire CI soon)
- ✓ Do not allow bypassing the above settings
- Save.
- Verify

```
git remote -v      # should show origin → github.com/<ORG/REPO>  
git branch -vv     # should show * main tracking origin/main
```



# Project Layout

```
apps/  
  web/          # Next.js (App Router)  
  api/          # Fastify/Express (REST for yellow-books)  
libs/  
  contract/     # Zod schemas + TypeScript types shared by web/api  
  config/       # Env loader (zod-validated), no secrets  
prisma/         # schema.prisma + migrations/ + seed script  
k8s/            # Kubernetes manifests (deployments, services, ingress)  
.github/workflows/ # GitHub Actions (CI/CD)
```

# Minimum required files

- `apps/web/package.json` → build script produces Next.js standalone output.
- `apps/api/package.json` → build compiles to `apps/api/dist`; start runs `node dist/main.js`.
- `libs/contract/index.ts` exports `YellowBookEntrySchema` + `YellowBookEntry` type.
- `prisma/schema.prisma` has `YellowBookEntry` model; `prisma/seed.ts` seeds  $\geq 5$  records.
- `k8s/` contains:
  - `api-deploy.yaml`, `web-deploy.yaml` (Deployments + Services)
  - `ingress-http.yaml` and `ingress-https.yaml` (ALB + optional TLS)
  - `job-migrate.yaml` (Prisma migrate deploy)
  - `hpa.yaml` (API autoscaling)
- `.github/workflows/cicd.yml` (build, push, deploy).

# Helpful root scripts (optional but nice)

```
// package.json (root)
{
  "scripts": {
    "build:web": "pnpm -C apps/web build",
    "build:api": "pnpm -C apps/api build",
    "lint": "nx run-many -t lint -p web api",
    "typecheck": "tsc -b --noEmit"
  }
}
```

- Separation of concerns: web vs api code paths are clean → easier Dockerfiles.
- Shared contracts: one source of truth for data shapes → fewer runtime bugs.
- Infra as code: k8s/ checked in → reproducible deployments.
- CI discoverability: workflows live in standard location so GitHub auto-detects.



# .dockerignore & Build Context

- Why it matters: Docker sends your build context (the files under .) to the daemon.
- Big contexts → slow builds, cache misses, higher CI cost.
- Add this at repo root:

```
# node & package managers
node_modules
pnpm-lock.yaml#ignore? (NO – keep for reproducible installs)

# VCS & artifacts
.git
*.log
coverage

# Next.js & builds
.next
out
dist

# tooling & caches
pnpm-store
.cache

# tests & storybook outputs (optional)
cypress
storybook-static

# env files (never commit secrets)
.env
.env.*
```



Хэрэглээний шинжлэх ухаан, Инженерчлэлийн сургууль  
Мэдээлэл, компьютерийн ухааны тэнхим

# Module 1 — Docker & ECR + GitHub Actions (Build & Push Now)



# Docker 101 (Image vs Container · Deep Dive)

Why Docker? Package once → run anywhere (your laptop, CI, Kubernetes).

- Image = read-only blueprint (layers). Built from a Dockerfile.
- Container = running instance of an image (ephemeral).
- Layers & cache: Each RUN/COPY/ADD creates a cacheable layer → order your Dockerfile to maximize reuse.
- Multi-stage builds: Build in a full toolchain image → copy only compiled artifacts to a tiny runtime image.
- Base images: node:20-alpine (small), node:20-bookworm-slim (more libs, fewer native module issues), distroless (very small, advanced).
- OCI: Docker images follow the OCI standard → works with many registries (ECR, GHCR, Docker Hub).



# Dockerfile anatomy

- FROM base → WORKDIR → COPY → RUN (install/build) → ENV → EXPOSE → CMD or ENTRYPOINT.

## Common pitfalls

- Copying the entire repo into the runtime stage → huge images.
- Putting secrets in **ENV** during build → secrets leak into image history.

# 1. Choose base image

```
FROM node:18
```

# 2. Set working directory inside the container

```
WORKDIR /app
```

# 3. Copy files from your computer into the container

```
COPY package*.json ./
```

```
RUN npm install
```

# 4. Copy the rest of the code

```
COPY . .
```

# 5. Define how to start your app

```
CMD ["npm", "start"]
```

# 6. Expose the port (optional)

```
EXPOSE 3000
```

```
FROM node:20-bookworm-slim AS deps
WORKDIR /app
COPY package.json pnpm-lock.yaml ./
RUN corepack enable && corepack prepare pnpm@9 --activate && pnpm fetch
COPY . .
RUN pnpm i --offline && pnpm build
```

```
FROM node:20-bookworm-slim AS runtime
WORKDIR /app
COPY --from=deps /app/dist ./dist
RUN useradd -m -u 10001 nodeuser
USER nodeuser
EXPOSE 3000
CMD ["node","dist/main.js"]
```

```
# inventory & layers
docker images && docker ps -a
docker history yb-api:local # see which steps are big
# clean up
docker rm -f $(docker ps -aq) # remove all containers (careful!)
docker rmi $(docker images -q) # remove all images (careful!)
```

# Registries 101 & Amazon ECR (Tags · Scanning · Lifecycle · Cross-Region)

Registry = storage for versioned images. We use Amazon ECR (private, regional, IAM-secured).

- URI format: `<acct>.dkr.ecr.<region>.amazonaws.com/yellowbooks/api:<tag>`.

## Tagging strategy

- Immutable: commit SHA (e.g., `:3a1b2c4`).
- Channel: `:main`, `:staging`, `:pr-123` (optional extra tag).
- For releases: `:v1.2.0` in addition to SHA.
- Prefer pinning by digest or SHA in deploys.

## Security & hygiene

- Enable Scan on push (detect CVEs in base images).
- Add Lifecycle policy (e.g., keep last 20 tags) to control storage cost.
- By default, ECR images are encrypted at rest (KMS).



# Manual login & push

Keep ECR and EKS in the same region to avoid pull latency and permission surprises.

```
1  aws ecr get-login-password --region <REGION> | \  
2  docker login --username AWS --password-stdin <ACCT>.dkr.ecr.<REGION>.amazonaws.com  
3  docker tag yb-api:local <ACCT>.dkr.ecr.<REGION>.amazonaws.com/yellowbooks/api:test  
4  docker push <ACCT>.dkr.ecr.<REGION>.amazonaws.com/yellowbooks/api:test
```

```
1  {  
2      "rules": [  
3          {  
4              "rulePriority": 1,  
5              "description": "Keep last 20 images (any tag)",  
6              "selection": { "tagStatus": "any", "countType": "imageCountMoreThan",  
7                  "countNumber": 20 },  
8              "action": { "type": "expire" }  
9          }  
10 ]  
}
```

# API Dockerfile

## (Multi-Stage · Fast Rebuilds · Secure Runtime)

- Goals: reproducible deps, small image, non-root, fast CI.
- Key choices:  
bookworm-slim base;  
pnpm workspace; copy  
only compiled output.

```
# apps/api/Dockerfile
# syntax=docker/dockerfile:1.6
FROM node:20-bookworm-slim AS base
WORKDIR /app

# deps - cache on lockfile
FROM base AS deps
RUN corepack enable && corepack prepare pnpm@9 --activate
COPY pnpm-workspace.yaml package.json pnpm-lock.yaml ./
COPY apps/api/package.json apps/api/package.json
COPY libs/contract/package.json libs/contract/package.json
RUN pnpm fetch --prod
COPY . .
RUN pnpm -C apps/api install --offline --prod=false

# build - compile TS → JS
FROM deps AS build
RUN pnpm -C apps/api build # emits apps/api/dist

# runtime - minimal files + non-root
FROM node:20-bookworm-slim AS runtime
ENV NODE_ENV=production
WORKDIR /app
COPY --from=deps /app/apps/api/package.json ./package.json
RUN corepack enable && corepack prepare pnpm@9 --activate \
  && pnpm i --prod --filter ./
COPY --from=build /app/apps/api/dist ./dist
RUN useradd -m -u 10001 nodeuser
USER nodeuser
EXPOSE 3000
CMD ["node", "dist/main.js"]
```



# Web Dockerfile (Next.js Standalone · Build-time vs Runtime Env)

- Goals: small server image that serves SSR/ISR; keep secrets out of build.
- Next.js config: ensure next.config.js sets output: 'standalone'.

```
# apps/web/Dockerfile
# syntax=docker/dockerfile:1.6
FROM node:20-bookworm-slim AS base
WORKDIR /app

FROM base AS deps
RUN corepack enable && corepack prepare pnpm@9 --activate
COPY pnpm-workspace.yaml package.json pnpm-lock.yaml ./
COPY apps/web/package.json apps/web/package.json
COPY libs/contract/package.json libs/contract/package.json
RUN pnpm fetch --prod
COPY . .
RUN pnpm -C apps/web install --offline --prod=false

FROM deps AS build
ENV NEXT_TELEMETRY_DISABLED=1
# Only NEXT_PUBLIC_* is safe at build - server secrets must be runtime
RUN pnpm -C apps/web build # emits .next/standalone + .next/static

FROM node:20-bookworm-slim AS runtime
ENV NODE_ENV=production PORT=3000 HOSTNAME=0.0.0.0
WORKDIR /app
COPY --from=build /app/apps/web/.next/standalone ./
COPY --from=build /app/apps/web/public ./public
COPY --from=build /app/apps/web/.next/static ./next/static
RUN useradd -m -u 10002 nextuser
USER nextuser
EXPOSE 3000
CMD ["node", "server.js"]
```



# Local Sanity Checklist (Run Images Before AWS)

## Typical issues & quick fixes

- API 500 on start →

wrong/missing

DATABASE\_URL; DB not reachable.

- Web cannot load data → wrong NEXT\_PUBLIC\_API (use http://localhost:3000).
- Port already in use → stop other processes or change -p mapping.
- CORS (if browser calls API) → allow origin http://localhost:3001 in API for local test.

## Success checklist

- /health returns 200 < 100ms.
- /yellow-books shows  $\geq 5$  records.
- docker history shows sane size (< ~300MB for class is fine).

```
export DATABASE_URL=postgres://user:pass@localhost:5432/yellow
```

```
docker build -f apps/api/Dockerfile -t yb-api:local .  
docker run --rm -p 3000:3000 -e DATABASE_URL yb-api:local  
# new terminal  
curl -s http://localhost:3000/health && echo
```

```
docker build -f apps/web/Dockerfile -t yb-web:local .  
docker run --rm -p 3001:3000 -e NEXT_PUBLIC_API=http://localhost:3000  
open http://localhost:3001/yellow-books
```

# Why Move CI Early? (Aha Now, Details Later)

## Pros

- Instant feedback loop: push  $\Rightarrow$  build  $\Rightarrow$  images in ECR.
- Mirrors real teams: CI produces artifacts on every commit.
- Motivates proper Dockerfiles and reproducible builds.

## Cons

- YAML, secrets, permissions can be overwhelming  $\rightarrow$  we limit scope to ECR push only this week.
- OIDC and cluster permissions not covered yet  $\rightarrow$  we use a temporary IAM user only for ECR (delete next week).



# Add GitHub Secrets & Variables (UI Walkthrough)

Where: GitHub → Repo → Settings → Secrets and variables → Actions.

Add these:

- Secrets: AWS\_ACCESS\_KEY\_ID, AWS\_SECRET\_ACCESS\_KEY.
- Variables: AWS\_REGION,  
ECR\_REPO=<acct>.dkr.ecr.<region>.amazonaws.com/yellowbooks,  
DOCKER\_BUILDKIT=1.

Tips:

- Use repo-level secrets for class; org-level if many repos.
- Masked by default; rotate if exposed.
- You can create environments (staging/prod) later with separate secrets.



# Actions Workflow (Build & Push Only — No Deploy Yet, with Caching)

- .github/workflows/ci-build-push.yml (annotated):
- [https://drive.google.com/file/d/1cFS2QIS6ubozD6O8jMcQtE-Drj\\_47mwo/view?usp=sharing](https://drive.google.com/file/d/1cFS2QIS6ubozD6O8jMcQtE-Drj_47mwo/view?usp=sharing)

## What this does

- Restores pnpm cache → faster builds.
- Uses buildx with GitHub cache for Docker layers.
- Pushes SHA tag every time; on main, also pushes :main tag.

# Nice-to-Have: Status Badge & PR Checks

Badge ([README.md](#)):

![CI](<https://github.com/<ORG/REPO>/actions/workflows/ci-build-push.yml/badge.svg>)

Require check on PRs:

- Settings → Branches → Branch protection → Require status checks to pass → select ci-build-push.

# Nice-to-Have: Status Badge & PR Checks

Badge ([README.md](#)):

![CI](<https://github.com/<ORG/REPO>/actions/workflows/ci-build-push.yml/badge.svg>)

Require check on PRs:

- Settings → Branches → Branch protection → Require status checks to pass → select ci-build-push.



# Homework & Rubric & Deliverables

- **Deliverables:** repo link, CI run link (green), ECR screenshots (both images with :<sha>), updated README badge.
- **Rubric (100):** Dockerfiles 30 · Local sanity 10 · ECR repos+policies 20 · CI build/push 30 · Docs 10.
- **Bonus (+10 means 1 point):** matrix build for push and pull\_request.

АСУУЛТ ?